TestPilot — Independent Review & Experiment Evidence Packet

Source: Updated user-provided November 2026 Word document. This is a searchable, Git-versioned textual transcription of the report; a separate build step publishes Word and PDF renderings. Original evidence labels are preserved. **Not an approved or signed academic finding.**

BIRLA INSTITUTE OF TECHNOLOGY & SCIENCE, PILANI

Work Integrated Learning Programmes Division

BSc Design and Computing · BSDCH ZC229T

TESTPILOT

Final Evaluation & Human Validation Packet

Prepared 2 November 2026 · EVIDENCE COLLECTION COPY — NOT A SIGNED RESEARCH RESULT

Purpose. Consolidates verifiable implementation evidence and the actual, documented 26 October Gemini comparison, provides a defensible preliminary AI assessment, and supplies ready-to-fill human rating and usability forms. No person, rating, experiment run, or session is fabricated.

1. Verified project evidence

Indicator Verified result Source/interpretation

Existing healthcare React/Node/MongoDB Healthcare app is tested, not built from

Application under test

deployment scratch for this project

Requirements 25 documented stories across 15 use cases Repository story catalog

Manual test-design baseline 36 cases Frozen manual case dataset

12 genuine candidates; 4 for each of 3 26 October archived AI comparison; NOT

Historical Gemini suggestions

stories 10 prompt iterations

29 Playwright PASS; 0 FAIL; 1 intentional

Latest automated CI GitHub Actions 41327612679; main branch

SKIP

Synthetic fixtures, not human-reviewed

Isolated recipe verification 6 scenario checks passed

study output

Render status: live, updated 1 November

Live Lab Deployed commit 042e14ae

2026

2. Preliminary AI versus manual design evaluation

The following is a provisional assistant assessment from the previously saved 26 October study. It must be confirmed or corrected by real independent reviewers. “Covered” means planned test steps and their expected outcomes sufficiently specify the criterion, not that the test was executed.

Preliminary manual

User story AI candidates Manual designs Preliminary AI criteria

criteria

US-001 Registration 4 6 2/2 2/2

US-005 Doctor browsing 4 3 0/2 2/2

US-017 Booking 4 4 2/2 2/2

Total 12 13 4/6 6/6

Interpretation: the manual baseline specifies all six sampled acceptance criteria, while the original Gemini candidate descriptions miss or inadequately assert two doctor-directory criteria. The sample is small and purposive; this does not show that manual design is universally better. There are no human time-savings or statistical significance estimates.

3. Specific evidence reviewers should inspect

US-001 — Register as patient

Two formal criteria: unique valid registration creates account/JWT; missing fields or short password are rejected. The archived Gemini case descriptions cover these at a coarse criterion level. Review the numeric mobile input,

seven/eight-character threshold, duplicate email normalization and explicit expected HTTP/state assertions. Compare against six manual cases.

US-005 — Browse approved doctors

Two criteria: only approved doctors appear; an empty directory is handled clearly. The original Gemini “approved-only” case checks response status and array presence but not that pending doctors are excluded. Its empty-list case lacks a precise empty-array/UI oracle. Compare with TP-011/TP-012 and add explicit assertions before approval.

US-017 — Book an appointment

Two criteria: valid booking creates Pending; unavailable slot is rejected. Four archived Gemini cases include success AI-001, occupied slot AI-002, missing fields AI-003 and no-auth AI-004. Manual TP-024 checks outside working sessions, TP-025 tests simultaneous booking and TP-026 checks conflict. Original model data contain schematic doctorId DOC-123 and slotTime rather than verified app payload; adapt ONLY in an isolated synthetic fixture. Original exported candidates remain NOT_RUN on production.

Reviewer rule: cite original case IDs and individual actions/expected results. A generic “works correctly” comment is insufficient. Record disagreements; do not silently average them.

4. Ten actual prompt runs — scoring register

These ten rows are a pre-registered PLAN, not evidence of completed model calls. Run them through the deployed Lab; preserve genuine model output and timestamp. A real reviewer must enter numeric scores 1–5 and note a specific observation for EACH iteration.

Genuine output Human C/A/H ratings 1– Reviewer note /

Run / Story Prompt variant

(file/time) 5 reference

01 US-001 Baseline _____ ___ / ___ / ___ _____

02 US-001 Boundary _____ ___ / ___ / ___ _____

03 US-001 Assertions _____ ___ / ___ / ___ _____

04 US-005 Baseline _____ ___ / ___ / ___ _____

05 US-005 Boundary _____ ___ / ___ / ___ _____

06 US-005 Assertions _____ ___ / ___ / ___ _____

07 US-017 Baseline _____ ___ / ___ / ___ _____

08 US-017 Boundary _____ ___ / ___ / ___ _____

09 US-017 Authorization _____ ___ / ___ / ___ _____

10 US-017 Assertions _____ ___ / ___ / ___ _____

C = completeness; A = accuracy; H = healthcare relevance. 1 = poor, 2 = weak, 3 = adequate, 4 = strong, 5 = excellent.

Evidence must be explicit (for example, missing exact 8-character assertion, invalid JSON field, or omitted concurrency case). Never record a zero as an unreviewed score; leave blank until rated.

Study collection link: https://testpilot-lab.onrender.com/ | Export name: testpilot-real-prompt-iterations.json. The Lab generation limit may interrupt consecutive runs; only resumed genuine provider responses count. Do not paste the private Lab access code into this document.

BITS Pilani · BSDCH ZC229T · TestPilot | Academic evidence — verify before submission

5. Independent reviewer sign-off sheets

Two different real QA reviewers each assess the SAME unchanged generated case set for all three stories. They must not copy each other's ratings. Enter "Covered / Not covered" for both AI and manual cases, cite relevant case

IDs/steps, then compare results and resolve each conflict.

Reviewer A — independent assessment

Reviewer initials: _____ Date/time (IST): _____ Original AI export/hash: _____

Criterion AI covered? Manual covered? Cited case ID / step / reason

US-001 / C1 _____ _____ _____

US-001 / C2 _____ _____ _____

US-005 / C1 _____ _____ _____

US-005 / C2 _____ _____ _____

US-017 / C1 _____ _____ _____

US-017 / C2 _____ _____ _____

Specific discrepancies / missing cases:

_____

I confirm that these judgments reflect my own inspection of the original AI and manual cases. Initials/signature:

_____

Reviewer B — independent assessment

Reviewer initials: _____ Date/time (IST): _____ Original AI export/hash: _____

Criterion AI covered? Manual covered? Cited case ID / step / reason

US-001 / C1 _____ _____ _____

US-001 / C2 _____ _____ _____

US-005 / C1 _____ _____ _____

US-005 / C2 _____ _____ _____

US-017 / C1 _____ _____ _____

US-017 / C2 _____ _____ _____

Specific discrepancies / missing cases:

_____

I confirm that these judgments reflect my own inspection of the original AI and manual cases. Initials/signature:

_____

Adjudication (after both reviewers submit)

Disagreements (criterion, A view, B view):

_____

Resolution with case-level evidence:

_____

Reviewer A sign-off: _____ Reviewer B sign-off: _____ Date: _____

BITS Pilani · BSDCH ZC229T · TestPilot | Academic evidence — verify before submission

6. Actual peer usability sessions

Two consenting real people must use TestPilot. An observer records actual task outcome and difficulty. Automated

browser tests and AI-authored imaginary feedback do not qualify as peer sessions.

Participant 1 — observer record

Participant initials/pseudonym: _____ Date: _____ Device/browser: _____

Consent to collect de-identified usability observations: YES / NO (If NO, do not record or submit.)

Actual start time: _____ Actual end time: _____ Measured duration (seconds): _____

Observed task Completed / Blocked / Not attempted Observed problem or time

Select and enter a documented story _____ _____

Generate and inspect Gemini suggestions _____ _____

Edit a case and identify a missing scenario _____ _____

Approve one case and reject another _____ _____

Interpret PASS versus NOT_RUN and

_____ _____

inspect evidence

Export reviewed cases / study log _____ _____

Participant’s actual feedback (quote or faithful notes):

_____

Observer’s improvement recommendation:

_____

Participant consent record reference: _____ Observer initials: _____

Participant 2 — observer record

Participant initials/pseudonym: _____ Date: _____ Device/browser: _____

Consent to collect de-identified usability observations: YES / NO (If NO, do not record or submit.)

Actual start time: _____ Actual end time: _____ Measured duration (seconds): _____

Observed task Completed / Blocked / Not attempted Observed problem or time

Select and enter a documented story _____ _____

Generate and inspect Gemini suggestions _____ _____

Edit a case and identify a missing scenario _____ _____

Approve one case and reject another _____ _____

Interpret PASS versus NOT_RUN and

_____ _____

inspect evidence

Export reviewed cases / study log _____ _____

Participant’s actual feedback (quote or faithful notes):

_____

Observer’s improvement recommendation:

_____

Participant consent record reference: _____ Observer initials: _____

BITS Pilani · BSDCH ZC229T · TestPilot | Academic evidence — verify before submission

7. Faculty validation and submission gate

The following are requirements for the final academic submission. The codebase and a draft report do NOT establish that these activities were performed. Mark complete only against actual evidence files.

Requirement Attach authentic evidence Verified now?

10 individual complete JSON responses /

Ten real Gemini provider runs NO

timestamps

10 reviewer judgments, rationales, review

Human 1–5 quality ratings NO

dates

US-001, US-005, US-017 records from

Two independent reviews per story NO

Reviewer A and B

Dispute resolution Signed rationale for disagreements NO

Measured tasks, independent feedback

Two consenting usability sessions NO

and consent

GitHub Actions 41327612679; 29 passed; 1

Technical execution evidence YES

skip

Final signed faculty report Updated conclusions + mentor approval NO

Final mentor: Kavi Arasu V Faculty confirmation/date: _____ Team confirmation/date:

_____

Research integrity note. The assistant may prepare preliminary case analysis and scoring instructions but is not an

independent human participant. Neither Playwright fixtures nor made-up human names, times, consent, or quotes should be submitted as observed findings.

Traceable sources: GitHub repository https://github.com/Akash-annavarapu/doctor-appointment-booking-system ;

verified main run https://github.com/Akash-annavarapu/doctor-appointment-booking-system/actions/runs/

41327612679 ; 26 October archival comparison TestPilot_AI_Comparison_2026-10-26.md ; user-stories.json; original

US-017 Gemini case JSON.

BITS Pilani · BSDCH ZC229T · TestPilot | Academic evidence — verify before submission